

Matrix Multiplication

CSCI 4576/5576

Outline

Memory and Matrix Multiplication

- Storage
- Blocking
- Cachegrind

Library examples

- HDF5
- MKL

Discuss homework

Warm-up: Memory

Program makes a memory reference

- If it's in cache, it gets returned immediately.
- If not: cache miss
 - New cache line is fetched

Cache Lines


- Good performance: use all the values in the line
- Bad performance: use a single element
 - Lots of memory bandwidth

Matrix storage

Matrix is 1D in memory


Row major

- $A(i,j) = A + i*n + j$



Column major

- $A(i,j) = A + i + j*n$



Example

```
int *M = new int [N*N];

// Access the matrix
for(int r=0; r<N; ++r)
{
    for(int c=0; c<N; ++c)
    {
        M[r*N + c] = r*c;
    }
}


// Cleanup
delete[] M;
```

Example

```
int *M = new int [N*N];

// Access the matrix
for(int r=0; r<N; ++r)
{
    for(int c=0; c<N; ++c)
    {
        M[r*N + c] = r*c;
    }
}

// Cleanup
delete[] M;
```






Matrix Multiply



```
for(int i=0; i<N; ++i)
{
    for(int j=0; j<N; ++j)
    {
        double sum = 0;
        for(int k=0; k<N; ++k)
        {
            sum = sum + A[i*N + k]*B[k*N + j];
        }
        M[i*N + j] = sum;
    }
}
```

Matrix Multiply


```
for(int i=0; i<N; ++i)
{
    for(int j=0; j<N; ++j)
    {
        double sum = 0;
        for(int k=0; k<N; ++k)
        {
            sum = sum + A[i*N + k]*B[k*N + j];
        }
        M[i*N + j] = sum;
    }
}
```



Block



Memory and Cachegrind



I1

D1

Ii

IId

example: row, col

```
int N = atoi(argv[1]);

// Allocate
int **M = new int * [N];
for (int i = 0; i < N; i++)
    M[i] = new int [N];

// Access the matrix
for(int r=0; r<N; ++r)
    for(int c=0; c<N; ++c)
        M[r][c] = r*c;

// Cleanup
for (int i = 0; i < N; i++)
    delete[] M[i];
delete[] M;
```

Cachegrind

			read	write
==91220==	I	refs:	433,595,222	
==91220==	I1	misses:	2,590	
==91220==	LLi	misses:	2,141	
==91220==	I1	miss rate:	0.00%	
==91220==	LLi	miss rate:	0.00%	
==91220==				
==91220==	D	refs:	277,732,903 (226,908,950 rd + 50,823,953 wr)	
==91220==	D1	misses:	1,589,774 (13,180 rd + 1,576,594 wr)	
==91220==	LLd	misses:	1,585,912 (9,563 rd + 1,576,349 wr)	
==91220==	D1	miss rate:	0.5% (0.0% + 3.1%)	
==91220==	LLd	miss rate:	0.5% (0.0% + 3.1%)	
==91220==				
==91220==	LL	refs:	1,592,364 (15,770 rd + 1,576,594 wr)	
==91220==	LL	misses:	1,588,053 (11,704 rd + 1,576,349 wr)	
==91220==	LL	miss rate:	0.2% (0.0% + 3.1%)	

example: col, row

```
int N = atoi(argv[1]);

// Allocate
int **M = new int * [N];
for (int i = 0; i < N; i++)
    M[i] = new int [N];

// Access the matrix
for(int c=0; c<N; ++c)
    for(int r=0; r<N; ++r)
        M[r][c] = r*c;

// Cleanup
for (int i = 0; i < N; i++)
    delete[] M[i];
delete[] M;
```

Cachegrind

```
==91221== I    refs:      433,595,222
==91221== I1   misses:      2,590
==91221== LLi  misses:      2,141
==91221== I1   miss rate:    0.00%
==91221== LLi  miss rate:    0.00%
==91221==
==91221== D    refs:      277,732,903 (226,908,950 rd + 50,823,953 wr)
==91221== D1   misses:      28,149,149 ( 3,137,555 rd + 25,011,594 wr)
==91221== LLd  misses:      23,167,621 (   59,194 rd + 23,108,427 wr)
==91221== D1   miss rate:    10.1% (   1.3% + 49.2% )
==91221== LLd  miss rate:     8.3% (   0.0% + 45.4% )
==91221==
==91221== LL   refs:      28,151,739 ( 3,140,145 rd + 25,011,594 wr)
==91221== LL   misses:      23,169,762 (   61,335 rd + 23,108,427 wr)
==91221== LL   miss rate:     3.2% (   0.0% + 45.4% )
```

