



Research Computing  
UNIVERSITY OF COLORADO

# Topologies and Communicators



# Outline

## Motivation

- Fox's algorithm

- Cannon's algorithm

## Topologies

- Structure imposed on the processes in a communicator that allows the processes to be addressed in different ways

## Communicators

- Collection of processes that can send messages to each other

# Topologies

# Topologies

Simplifies code

Can allow MPI to optimize communication

Creating a new topology provides:

- A new communicator

- Mapping functions: rank  $\rightarrow$  topology  $\rightarrow$  rank

## Types

- Cartesian: processes can be identified by coordinates

- Graph

# Creating a Cartesian Topology

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims, int *dims,  
                   int *periods, int reorder, MPI_Comm *comm_cart);
```

comm\_old: existing communicator

ndims: number of dimensions

\*dims: the actual dimension

\*periods: logical array indicating whether a dimension is cyclic

reorder: logical, false = preserve rank order

comm\_cart: new communicator

```
int MPI_Dims_create(int nnodes, int ndims, int *dims);
```

Creates a division of processors in a cartesian grid

# Example

```
const int dim = 2;  
int grid[dim] = {0,0};  
  
// Assign the grid dimensions  
MPI_Dims_create(size, dim, grid);  
  
// The new communicator  
MPI_Comm comm_grid;  
  
// Allow cyclic behavior  
int preiodic[dim] = {TRUE,TRUE};  
  
// Create the communicator  
MPI_Cart_create(comm, dim, grid, preiodic, TRUE, &comm_grid);
```

# Mapping functions

```
int MPI_Cart_coords(MPI_Comm comm, int rank, int maxdims, int *coords);
```

Determines process coords in cartesian topology given rank in group

```
int MPI_Cart_rank(MPI_Comm comm, int *coords, int *rank);
```


Determines process rank in communicator given Cartesian location

```
int MPI_Cart_shift(MPI_Comm comm, int direction, int displ,  
                  int *source, int *dest);
```


Returns source and destination ranks, given a shift direction and amount

direction = 0, direction=1, ... direction=dim-1

displ > 0



displ < 0



# Example

// What is the rank grid coordinate?

```
int grid_coord[dim] = {0,0};
```

```
MPI_Cart_coords(comm_grid, rank, dim, grid_coord);
```

// Given a coordinate, what is the rank?

```
int rank_id;
```

```
MPI_Cart_rank(comm_grid, grid_coord, &rank_id);
```

// Who is my neighbor to the right?

```
int neighbor;
```

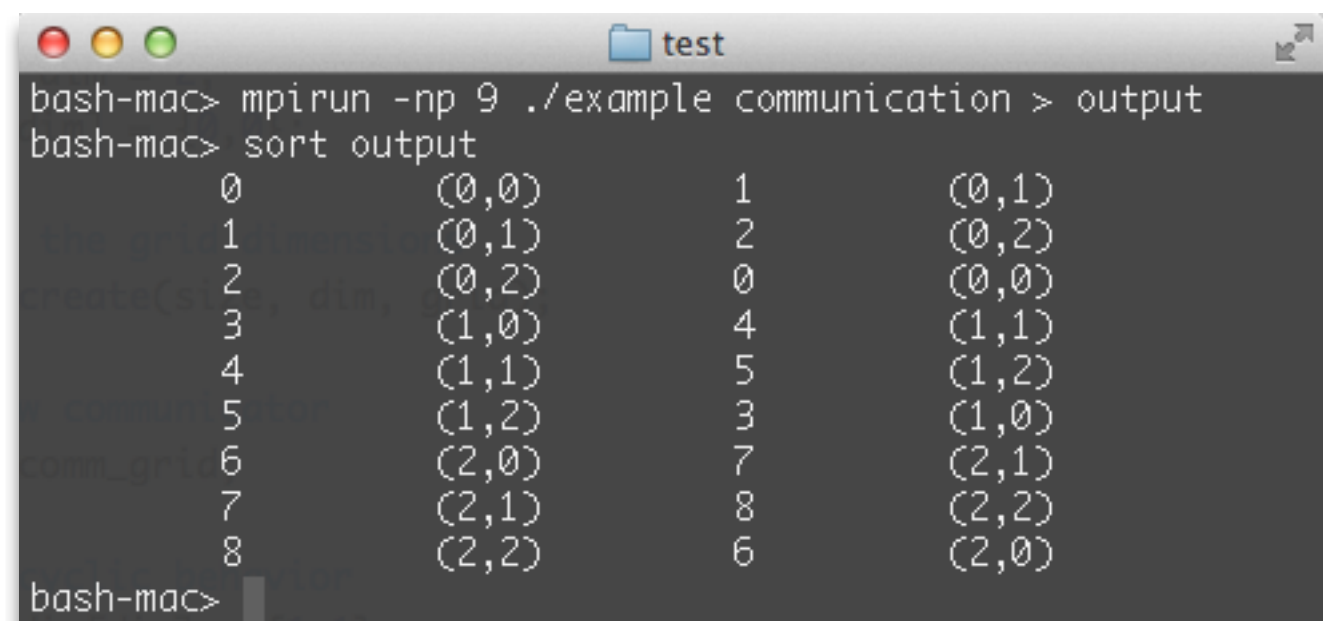
```
MPI_Cart_shift(comm_grid, 1, 1, &rank, &neighbor);
```

```
int neighbor_coord[dim] = {0,0};
```

```
MPI_Cart_coords(comm_grid, neighbor, dim, neighbor_coord);
```

// print rank\_id grid\_coord

// neighbor, neighbor\_coord



A terminal window titled 'test' showing the execution of an MPI program. The command 'bash-mac> mpirun -np 9 ./example communication > output' is entered, followed by 'bash-mac> sort output'. The output is a 3x3 grid of rank IDs and their corresponding grid coordinates (rank, (x,y)).

|   |       |   |       |
|---|-------|---|-------|
| 0 | (0,0) | 1 | (0,1) |
| 1 | (0,1) | 2 | (0,2) |
| 2 | (0,2) | 0 | (0,0) |
| 3 | (1,0) | 4 | (1,1) |
| 4 | (1,1) | 5 | (1,2) |
| 5 | (1,2) | 3 | (1,0) |
| 6 | (2,0) | 7 | (2,1) |
| 7 | (2,1) | 8 | (2,2) |
| 8 | (2,2) | 6 | (2,0) |

The terminal ends with 'bash-mac>' and a cursor.



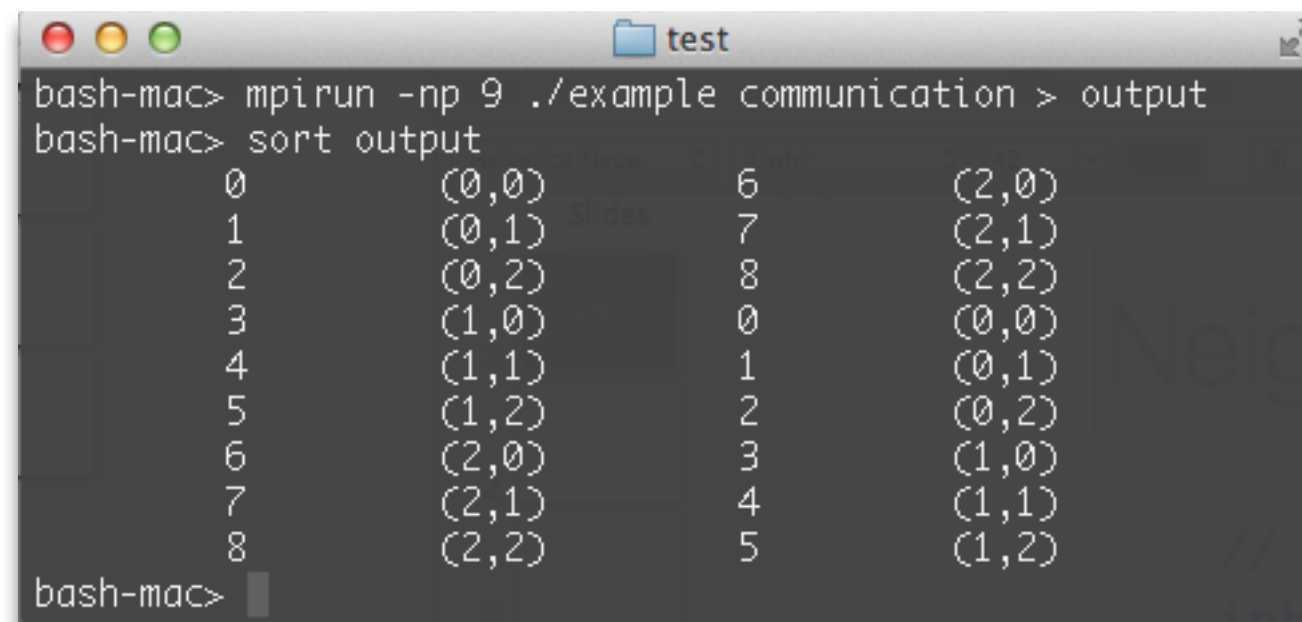
# Neighbor above?

```
// Who is my neighbor to the right?
int neighbor;
MPI_Cart_shift(comm_grid, 1, 1, &rank, &neighbor);

int neighbor_coord[dim] = {0,0};
MPI_Cart_coords(comm_grid, neighbor, dim, neighbor_coord);

// Who is my neighbor above?
int neighbor;
MPI_Cart_shift(comm_grid, 0, -1, &rank, &neighbor);

int neighbor_coord[dim] = {0,0};
MPI_Cart_coords(comm_grid, neighbor, dim, neighbor_coord);
```

A terminal window titled 'test' showing the execution of an MPI program. The user runs 'bash-mac> mpirun -np 9 ./example communication > output' and then 'bash-mac> sort output'. The output is a 4x4 grid of coordinates (x,y) for ranks 0 through 8, arranged in a 3x3 grid pattern. The first three ranks (0,1,2) are in the first row, the next three (3,4,5) in the second, and the last three (6,7,8) in the third. The coordinates are (0,0), (0,1), (0,2) for the first row, (1,0), (1,1), (1,2) for the second, and (2,0), (2,1), (2,2) for the third.

```
bash-mac> mpirun -np 9 ./example communication > output
bash-mac> sort output
0      (0,0)      6      (2,0)
1      (0,1)      7      (2,1)
2      (0,2)      8      (2,2)
3      (1,0)      0      (0,0)
4      (1,1)      1      (0,1)
5      (1,2)      2      (0,2)
6      (2,0)      3      (1,0)
7      (2,1)      4      (1,1)
8      (2,2)      5      (1,2)
bash-mac>
```

# Communicators

# Communicators

Group of processes that can send and receive messages

Perform collective communications

Creation:

Create groups of processors -> then create the MPI\_Comm

MPI\_Comm\_split: split an existing communicator

MPI\_Cart\_sub: create from a topology

When you are done using them...

MPI\_Comm\_free

# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm, int split_key, int key,  
                  MPI_Comm *newcomm);
```

split\_key: processes with the same split\_key will be in the same communicator

key: control of rank assignment

| rank_id | grid_coord[0] | grid_coord[1] |
|---------|---------------|---------------|
| 0       | 0             | 0             |
| 1       | 0             | 1             |
| 2       | 1             | 0             |
| 3       | 1             | 1             |

|   |   |
|---|---|
| 0 | 1 |
| 2 | 3 |

```
MPI_Comm comm_row;  
MPI_Comm_split(comm_grid,                                &comm_row);  
MPI_Comm_split(comm_grid, grid_coord[0], rank_id, &comm_row);
```

# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm, int *remain_dims,  
                 MPI_Comm *comm_new);
```

Partitions a communicator into subgroups which form lower-dimensional cartesian subgrids

remain\_dims: ith dimension is kept in the subgrid (true) or is dropped (false)

```
// Col communicator  
MPI_Comm comm_col;
```

```
// What goes here? (true, false)  
int free_coords[dim] =
```

```
// Creat the communicator  
MPI_Cart_sub(comm_grid, free_coords, &comm_col);
```

# Homework 6



# Tasks

## File I/O

No parallel IO, you can use existing code

Partial read and distribute

For best performance, you will need to “chunk”

## Matrix Multiply

General Comments

Cannon's Algorithm

Fox's Algorithm

## Advice

# Example


Only blocks of the same color can be multiplied










# Cannon's Algorithm

Rearrange the read and distribute




$$p(i, j) \rightarrow A(i, k)B(k, j)$$











Step 1









Step 2



Step 3



Step 4



# Fox's Algorithm

Read in normally





# Step 1

Broadcast diagonal element to each member of row



Multiply



# Step 2



Shift, Multiply





Step 3



Shift, Multiply



Step 4



Shift, Multiply



# Done



# Advice

Start early

Don't worry about IO at first.

- Read in entire matrix

- Send the correct components

Consider a class or struct for keeping information

Use a simple derived MPI\_Datatype

- e.g. 1D array for matrix representation