



Research Computing
UNIVERSITY OF COLORADO

MPI Datatypes

Creating Custom Datatypes

What is a custom MPI datatype?

Construction

`MPI_Type_contiguous`

`MPI_Type_vector`

`MPI_Type_struct`

Commit

`MPI_Type_commit`

Send and receive

Free

`MPI_Type_free`

Point to point communication



```
int MPI_Send( void *buf,  
              int count,  
              MPI_Datatype dt,  
              int dest,  
              int tag,  
              MPI_Comm comm );
```

```
int MPI_Recv( void *buf,  
              int count,  
              MPI_Datatype dt,  
              int source,  
              int tag,  
              MPI_Comm comm,  
              MPI_Status *status );
```

`int source`


envelope

`int dest`

`int tag`

`MPI_Comm comm`

message body



Communication review

Consider the following code:

```
int data = 10;
if (rank == 0)
    MPI_Send(&data, 1, MPI_INT, 1, tag, comm);
else
{
    MPI_Recv(&data, 1, MPI_INT, 0, tag, comm, &status);
    std::cout << "Message received!" << std::endl;
}
```

envelope

message body

If this code is run on 1 processor, what do you expect to happen?

error condition

If the code is run on 3 processors, what do you expect?

code will hang

MPI Datatypes

```
MPI_Send(&data, count, MPI_INT, 1, tag, comm);
```



What do we need to know?

address, datatype, number of elements


What does MPI do?

parses memory based on the starting address, size, and count


based on *contiguous* data



MPI Custom Datatypes



Where do the types start?



How many of each type?

What type of data?



count[0] = 1

count[1] = 1

count[2] = 2


types[0] = MPI_INT

types[1] = MPI_FLOAT

types[2] = MPI_DOUBLE

MPI Datatype is a map

```
struct {  
    int num;  
    float x;  
    double data[2];  
} obj;
```



MPI_Datatype obj_type;

this is what we are
learning today

type	disp	count
MPI_INT	address	1
MPI_FLOAT	address	1
MPI_DOUBLE	address	2

MPI_Bcast(&obj, 1, obj_type, 0, comm);

memory address

this is how you interpret the data

Contiguous

```
int MPI_Type_contiguous(int count,  
                        MPI_Datatype old_type,  
                        MPI_Datatype new_type);
```

Example

```
int B[2][3];  
MPI_Datatype matrix;  
MPI_Type_contiguous(6, MPI_INT, &matrix);
```

int	int	int	int	int	int
datatype	displacement	count			
MPI_INT	0	6			

What's the advantage over this?

```
MPI_Send(B, 6, MPI_INT, 1, tag, comm);
```


Contiguous Example

```
const int N = 10;

double A[N][N];
double B[N][N];

MPI_Datatype matrix;

MPI_Type_contiguous(N*N, MPI_DOUBLE, &matrix);
MPI_Type_commit(&matrix);

if (rank == master_rank
    MPI_Send(A, 1, matrix, 1, 10, comm);
else if( rank == 1 )
    MPI_Recv(B, 1, matrix, 0, 10, comm, &status);
```

Vector

User specifies memory locations

```
int MPI_Type_vector(int count,  
                    int blocklength,  
                    int stride,  
                    MPI_Datatype old_type,  
                    MPI_Datatype *newtype);
```

newtype has

count blocks each consisting of blocklength copies of oldtype.

Displacement between blocks is set by stride.

Example

```
const int N = 5;  
MPI_Datatype dt;  
MPI_Type_vector(N, 1, N, MPI_INT, &dt);
```

Variables

count = 5

blocklength = 1

stride = 5

How is this useful?



Matrix Column Example

```
const int N = 10;
```

```
double A[N][N];
```

```
MPI_Datatype column;
```

```
MPI_Type_vector(N, 1, N, MPI_DOUBLE, &column);
```

```
MPI_Type_commit(&column);
```

```
if (rank == master_rank)
```

```
    MPI_Send(&A[0][2], 1, column, 1, 10, comm);
```

```
else if( rank == 1 )
```

```
    MPI_Recv(&A[0][2], 1, column, 0, 10, comm, &status);
```

Extent

Memory span of a datatype

```
int MPI_Type_extent(MPI_Datatype datatype,  
                    MPI_Aint *extent);
```

Usage: think “size of”

```
MPI_Aint intex;
```

```
MPI_Type_extent(MPI_INT, &intex);
```

```
displacements[0] = static_cast<MPI_Aint>(0);  
displacements[1] = intex;
```


Structure

Heterogeneous

Most general derived datatype

```
int MPI_Type_struct(int count,  
                    int blocks[],  
                    MPI_Aint displacements[],  
                    MPI_Datatype types[],  
                    MPI_Datatype *newtype);
```

newtype consists of

count blocks where the *ith* block is blocks[i] copies of the type types[i]

The displacement of the *ith* block (in bytes) is given by displacements[i]

Example

```
struct {  
    int num;  
    float x;  
    double data[4];  
} obj;
```

Variables

count =

blocks =

types =

displacements =

count blocks where the *ith* block is

blocks[i] copies of the type types[i]

The displacement of the *ith* block (in bytes) is
given by displacements[i]



Example

```
int blocks[3]={1,1,4};
MPI_Datatype types[3]={MPI_INT, MPI_FLOAT, MPI_DOUBLE};
MPI_Aint displacements[3];

MPI_Datatype obj_type;
MPI_Aint intex, floatex;


MPI_Type_extent(MPI_INT, &intex);
MPI_Type_extent(MPI_FLOAT, &floatex);
displacements[0] = static_cast<MPI_Aint>(0);
displacements[1] = intex;
displacements[2] = intex+floatex;
MPI_Type_struct(3, blocks, displacements, types, &obj_type);
```

C/C++

```
struct {  
    int n;  
    double x[3];  
} obj;
```



```
struct {  
    int n;  
    double * x;  
} obj;
```



MPI_Address

```
int MPI_Address(void *location, MPI_Aint *address);
```

Gets the address location in memory

```
struct {  
    int n;  
    double * x;  
} obj;
```


Commit and Free

After construction, you must commit the datatype

```
int MPI_Type_commit(MPI_Datatype *datatype);
```

When you are done, you need to free the datatype

```
int MPI_Type_free(MPI_Datatype *datatype);
```

Send are Receive

```
if (rank==3){  
  
    obj.num=6;  
    obj.x=3.14;  
    for(int i=0;i<4;++i)  
        obj.data[i]=(double) i;  
  
    MPI_Send(&obj,1,obj_type,1,52,comm);  
  
} else if(rank==1) {  
  
    MPI_Recv(&obj,1,obj_type,3,52,comm,&status);  
  
}
```

Broadcast

```
    if (rank == master_rank)
    {
        for(int i=0; i<obj.n; ++i)
            obj.x[i] = i+1;
    }

    MPI_Bcast(&obj,1,obj_type,0,comm);
```